

# INFO 6205

## Program Structures & Algorithms

### Fall 2020

### Assignment No 5

- **Task**

The task given to us was to implement a parallel sorting algorithm by leveraging concurrent programming in Java. We use `CompletableFuture` provided by Java 8 which is a specific implementation of the `Future` interface which allows users to explicitly complete an asynchronous computation. In this assignment, we conduct various experiments to find a suitable value of the number of threads and the cutoff value for which we should stop our parallel computation and resort to system sort instead.

- **Implementation**

In `ParSort.java`, I implemented the algorithm for merging (or combining) the two parallel partition computations and provided a variable for explicitly setting the thread pool size by using `ForkJoinPool`. Essentially, in this parallel sorting algorithm, we are simply doing merge sort in parallel. I used `ForkJoinPool` which is a specific implementation of `Executor Service` which provides efficient processing by means of work stealing where threads in the pool attempt to execute tasks created by other threads in the pool. In `Main.java`, I have different methods to implement different stages of my experiment but the main itself simply runs the parallel sorting algorithm at different thread pool sizes ranging from 1 to 16 (increasing by a factor of 2) and different cutoffs (ranging from 500,000 to 1,000,000) for an array size of 2,000,000. I plot results for every stage of my experiment using `Plotter.java`.

Other stages of my experiment are conducted in methods:

1. `experimentWithFixedThreadSizeVaryingCutoffs()`: Choose a specific thread size and plot the cutoff vs time graph
2. `experimentWithVaryingThreadSizeFixedCutoffs()`: Choose a specific cutoff and plot the time vs size graph with varying thread sizes

- **Output**

I ran my parallel sorting algorithm with cutoffs ranging from 500,000 to 1,000,000 for array sizes: 1,000,000, 2,000,000 and 4,000,000 and thread pool sizes: 1, 2, 4, 8, 16 at 20 runs each. I documented the average time taken for each array at each thread pool size, the minimum time taken and maximum time taken, the corresponding cutoff at which that value was taken and plotted the Cutoff vs Time graph for each. The results are presented in Table 3.1, 3.2, 3.3 and the graphs are presented in Figure 3.1, 3.2, 3.3.

| Threads | Average | Minimum Time | Minimum Cutoff | Maximum Time | Maximum Cutoff |
|---------|---------|--------------|----------------|--------------|----------------|
| 1       | 104 ms  | 97 ms        | 600,000        | 129 ms       | 530,000        |
| 2       | 63 ms   | 59 ms        | 590,000        | 67 ms        | 510,000        |
| 4       | 66 ms   | 63 ms        | 630,000        | 74 ms        | 730,000        |
| 8       | 63 ms   | 59 ms        | 750,000        | 77 ms        | 520,000        |
| 16      | 63 ms   | 59 ms        | 560,000        | 77 ms        | 910,000        |

*Table 3.1: Results recording average time, minimum and maximum time and the corresponding cutoffs when the array size = 1,000,000*




Figure 3.1: Graph recording Cutoff vs Time for array size = 1,000,000 at 20 runs

| Threads | Average | Minimum Time | Minimum Cutoff | Maximum Time | Maximum Cutoff |
|---------|---------|--------------|----------------|--------------|----------------|
| 1       | 195 ms  | 186 ms       | 970,000        | 216 ms       | 700,000        |
| 2       | 185 ms  | 176 ms       | 970,000        | 198 ms       | 950,000        |
| 4       | 130 ms  | 125 ms       | 700,000        | 143 ms       | 580,000        |
| 8       | 110 ms  | 103 ms       | 830,000        | 123 ms       | 530,000        |
| 16      | 116 ms  | 102 ms       | 102,000        | 171 ms       | 800,000        |

Table 3.2: Results recording average time, minimum and maximum time and the corresponding cutoffs when the array size = 2,000,000




Figure 3.2: Graph recording Cutoff vs Time for array size = 2,000,000 at 20 runs

| Threads | Average | Minimum Time | Minimum Cutoff | Maximum Time | Maximum Cutoff |
|---------|---------|--------------|----------------|--------------|----------------|
| 1       | 383 ms  | 350 ms       | 910,000        | 538 ms       | 610,000        |
| 2       | 412 ms  | 360 ms       | 780,000        | 500 ms       | 840,000        |
| 4       | 352 ms  | 336 ms       | 880,000        | 408 ms       | 540,000        |
| 8       | 349 ms  | 310 ms       | 940,000        | 393 ms       | 550,000        |
| 16      | 291 ms  | 249 ms       | 660,000        | 508 ms       | 720,000        |

Table 2.3: Results recording average time, minimum and maximum time and the corresponding cutoffs when the array size = 4,000,000




Figure 3.3 Graph recording Cutoff vs Time for array size = 4,000,000 at 20 runs

### Console Output for Array Size = 4,000,000

```
cutoff and threads = 1: 510000      20 times Time:8053ms
cutoff and threads = 1: 520000      20 times Time:7514ms
cutoff and threads = 1: 530000      20 times Time:7374ms
cutoff and threads = 1: 540000      20 times Time:7590ms
cutoff and threads = 1: 550000      20 times Time:7573ms
cutoff and threads = 1: 560000      20 times Time:7212ms
cutoff and threads = 1: 570000      20 times Time:7292ms
cutoff and threads = 1: 580000      20 times Time:7216ms
cutoff and threads = 1: 590000      20 times Time:7507ms
...
...
Average time taken over all cutoffs for thread count: 1 = 383
Minimum time taken over all cutoffs for thread count: 1 = 350 at cutoff: 910000
Maximum time taken over all cutoffs for thread count: 1 = 538 at cutoff: 610000
...
...
cutoff and threads = 16: 950000      20 times Time:5143ms
cutoff and threads = 16: 960000      20 times Time:5094ms
cutoff and threads = 16: 970000      20 times Time:5278ms
cutoff and threads = 16: 980000      20 times Time:5221ms
cutoff and threads = 16: 990000      20 times Time:5565ms
cutoff and threads = 16: 1000000     20 times Time:5514ms
Average time taken over all cutoffs for thread count: 16 = 291
Minimum time taken over all cutoffs for thread count: 16 = 249 at cutoff: 660000
Maximum time taken over all cutoffs for thread count: 16 = 508 at cutoff: 720000
```

- **Conclusions and Supporting Evidence**

The idea behind finding the cutoff at which the time was minimum and maximum was to observe a pattern but from the above results we can see that there is no specific value at which we are to find a minimum or a maximum. Now, the question arises that is our cutoff range not big enough? To answer this question, I increased my range to test the time for cutoff sizes from 100,000 to 1,000,000 for a fixed array size of 1,000,000 at a fixed thread pool size of 16. I chose 16 because from the results above it was clear that the minimum time is always found at thread pool size of 16.




Figure 4.1 Graph recording time to do parallel sorting for cutoff range 100,000 – 1,000,000, fixed thread pool size = 16, and array size = 1,000,000 at 50 runs

From the graph above, we see that the time taken when the cutoff is 100,000 is very high, it slowly starts decreasing as our cutoff increases and more or less stabilizes to a consistent value from cutoff value = 500,000. However, in order to confirm if the time ever goes beyond our specific range, I decided to run the parallel sorting algorithm for a cutoff starting from 100, which as expected gave an out of memory error. So, I started the sort algorithm for a cutoff range: 1,000 to 591,000 with increments of 10,000 and plotted the results. I run the experiment for thread pool size = 1 and thread pool size = 16 as they gave the worst and best performance respectively.




Figure 4.2 Graph comparing time to do parallel sorting for cutoff range 1000 – 591,000 and fixed thread size = 16, fixed thread size = 1, and array size = 1,000,000 at 50 runs


We can see that the graph is initially inverse of what we got for the range 500,000 to 1,000,000 where Thread Pool Size = 1 performed the poorest and Thread Pool Size = 16 performed the best. Instead, here we see that for thread pool size = 1, we get better results initially as opposed to thread pool size = 16. As we slowly start moving towards the cutoff range > 100,000, our results altogether start improving, following the improvement with increase in thread pool size trend. We can credit this decrease in performance with an imbalance in the cost of splitting work, queue and job handling and aggregating jobs results. There comes a point in the parallel algorithms where the overhead from

waiting for threads to finish, creating and managing threads simply outweighs any performance improvements we achieve from concurrent computations. Hence, the cutoff value is extremely important from where the algorithm switches to sequential sorting. In fact, the time taken by system sort at array size = 1,000,000 was 105 ms which is lower than the time achieved by both thread pool = 1, thread pool = 16 at cutoff = 1,000.

Choosing cutoff = 1,000, running the parallel sorting algorithm for varying thread sizes and comparing it with the system sort we find results that confirm our hypothesis that it is significantly better to use normal sequential merge sort as the partition becomes smaller.

| Array Size | Thread Pool Size = 1 | Thread Pool Size = 2 | Thread Pool Size = 4 | Thread Pool Size = 8 | Thread Pool Size = 16 | System Sort |
|------------|----------------------|----------------------|----------------------|----------------------|-----------------------|-------------|
| 200,000    | 36 ms                | 42 ms                | 64 ms                | 48 ms                | 59 ms                 | 21 ms       |
| 400,000    | 55 ms                | 62 ms                | 51 ms                | 68 ms                | 94 ms                 | 40 ms       |
| 800,000    | 118 ms               | 126 ms               | 127 ms               | 127 ms               | 147 ms                | 83 ms       |
| 1,600,000  | 351 ms               | 446 ms               | 510 ms               | 511 ms               | 483 ms                | 194 ms      |
| 3,200,000  | 1464 ms              | 1654 ms              | 1151 ms              | 970 ms               | 984 ms                | 365 ms      |

*Table 4.1: Results recording time taken by parallel sorting algorithm at various thread pool sizes at cutoff = 1,000 and the system sort algorithm*



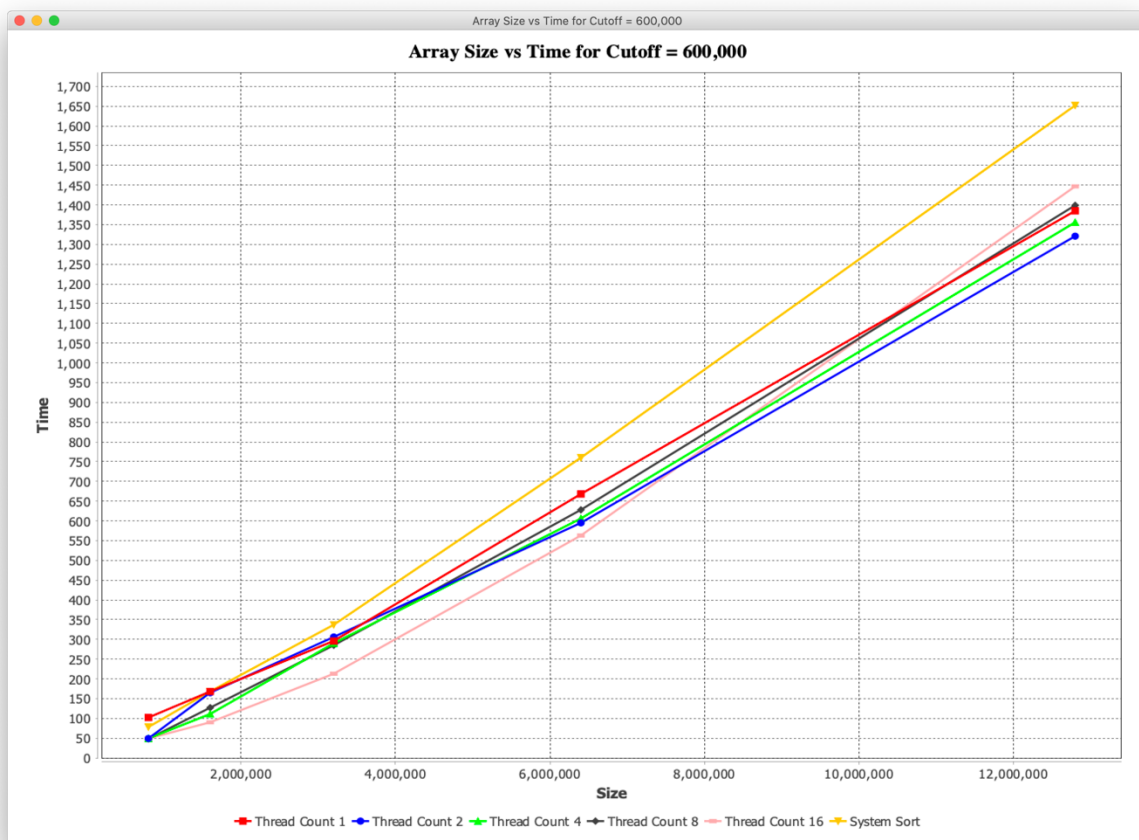
*Figure 4.3: Graph plotting Size vs Time graph for parallel sorting algorithm with different thread pool sizes at cutoff = 1,000 and the system sort*



Finally, choosing a good cutoff value = 600,000, I reported and plotted below the results of the performance of my parallel sorting algorithm with each thread pool sizes, system sort and array size ranging from 800,000 – 12,800,000.

| Array Size | Thread Pool Size = 1 | Thread Pool Size = 2 | Thread Pool Size = 4 | Thread Pool Size = 8 | Thread Pool Size = 16 | System Sort |
|------------|----------------------|----------------------|----------------------|----------------------|-----------------------|-------------|
| 800,000    | 102 ms               | 49 ms                | 49 ms                | 48 ms                | 49 ms                 | 78 ms       |
| 1,600,000  | 168 ms               | 165 ms               | 111 ms               | 127 ms               | 90 ms                 | 168 ms      |
| 3,200,000  | 296 ms               | 306 ms               | 290 ms               | 285 ms               | 213 ms                | 337 ms      |
| 6,400,000  | 668 ms               | 595 ms               | 606 ms               | 628 ms               | 563 ms                | 760 ms      |
| 12,800,000 | 1385 ms              | 1321 ms              | 1356 ms              | 1399 ms              | 1447 ms               | 1652 ms     |

*Table 4.2: Results recording time taken by parallel sorting algorithm at various thread pool sizes at cutoff = 600,000 and the system sort algorithm*



*Figure 4.4: Graph plotting Size vs Time graph for parallel sorting algorithm with different thread pool sizes at cutoff = 600,000 and the system sort*

#### **Console Output:**

For thread count: 1 and size: 800000 the time taken is: 102



```
For thread count: 2 and size: 800000 the time taken is: 49
For thread count: 4 and size: 800000 the time taken is: 49
For thread count: 8 and size: 800000 the time taken is: 48
For thread count: 16 and size: 800000 the time taken is: 49
Time taken by system sort is: 78
For thread count: 1 and size: 1600000 the time taken is: 168
For thread count: 2 and size: 1600000 the time taken is: 165
For thread count: 4 and size: 1600000 the time taken is: 111
For thread count: 8 and size: 1600000 the time taken is: 127
For thread count: 16 and size: 1600000 the time taken is: 90
Time taken by system sort is: 168
For thread count: 1 and size: 3200000 the time taken is: 296
For thread count: 2 and size: 3200000 the time taken is: 306
For thread count: 4 and size: 3200000 the time taken is: 290
For thread count: 8 and size: 3200000 the time taken is: 285
For thread count: 16 and size: 3200000 the time taken is: 213
Time taken by system sort is: 337
For thread count: 1 and size: 6400000 the time taken is: 668
For thread count: 2 and size: 6400000 the time taken is: 595
For thread count: 4 and size: 6400000 the time taken is: 606
For thread count: 8 and size: 6400000 the time taken is: 628
For thread count: 16 and size: 6400000 the time taken is: 563
Time taken by system sort is: 760
For thread count: 1 and size: 12800000 the time taken is: 1385
For thread count: 2 and size: 12800000 the time taken is: 1321
For thread count: 4 and size: 12800000 the time taken is: 1356
For thread count: 8 and size: 12800000 the time taken is: 1399
For thread count: 16 and size: 12800000 the time taken is: 1447
Time taken by system sort is: 1652
```

From the above results, we can see that our parallel algorithm almost follows the same linearithmic pattern as the system sort but the constant value changes according to the thread pool count. We can see that for a good cutoff value, we can benefit from using a parallel sorting algorithm with a bigger thread pool size.

## • Conclusions

Summarizing our observations from our experiments, we find that:

1. This experiment will differ majorly on different systems and may produce different results.
2. When the cutoff value lies between 100,000 – 1,000,000, our parallel sorting algorithm performs the best with a peak performance between 500,000 – 700,000 which may vary given different conditions. (Fig. 4.1)
3. In the cutoff values between 500,000 – 1,000,000, our parallel sorting algorithm performs the best when thread pool size = 16, and the worst when thread pool size = 1. (Table 3.1 – 3.3, Fig 3.1 – 3.3).
4. When the cutoff size is too small (most of our work is being done in parallel), we face the consequences of high overhead of creating, managing and synchronization and the parallel sorting algorithm becomes an overkill. The cost of these outweighs any improvement in performance, hence, we must be careful in the cutoff value we choose. (Fig 4.2, Fig 4.3)
5. When cutoff size is too small, sequential system sort (normal merge sort) performs better. (Fig 4.3)
6. The parallel sorting algorithm has an approximately linearithmic performance similar to the system sort algorithm, however, the constant changes with respect to the thread pool size being used. (Fig 4.4, Table 4.2).
7. For parallel programming, we must always perform experiments to provide a reasonable cutoff because there comes a point where the cost of creating, managing and synchronizing threads exceeds any performance improvement we could've gotten.